

```
hive> create database dsbda;
OK
Time taken: 2.201 seconds
hive> use dsbda;
OK
Time taken: 0.019 seconds
hive> create external table empdata( ename string, esal int)
    > row format delimited fields terminated by "," stored as textfile location
    > "/home/hduser/Desktop/empdata";
OK
Time taken: 0.148 seconds
hive> show tables;
OK
empdata
Time taken: 0.066 seconds, Fetched: 1 row(s)
hive> load data local inpath '/home/hduser/empdb.txt' into table empdata;
Loading data to table dsbda.empdata
OK
Time taken: 0.793 seconds
hive> select * from
    > empdata;
OK
sachin  5000
yuvraj  3000
virat   1000
rahul   4000
saurav  3500
scott   2500
smith   1500
mathew  2400
Laxman  4500
ricky   1300
ram     2000
Time taken: 0.528 seconds, Fetched: 11 row(s)
hive> alter table empdata change esal esal string;
OK
Time taken: 0.054 seconds
hive> describe empdata;
OK
ename                string
esal                  string
Time taken: 0.037 seconds, Fetched: 2 row(s)
hive> drop table empdata;
OK
hive> create table empinfo(empno int, empgrade string) row format delimited fields
terminated by ',' stored as textfile;
OK
Time taken: 0.045 seconds
hive> load data local inpath '/home/hduser/empinfo.txt' into table empinfo;
Loading data to table dsbda.empinfo
OK
Time taken: 0.141 seconds
hive> select * from empinfo;
OK
1      A
2      B
3      B
4      A
5      B
6      B
7      D
8      E
9      C
10     C
11     D
Time taken: 0.053 seconds, Fetched: 11 row(s)
```

```
hive> create table empdbnew(eno int, ename string, esal int) row format delimited fields
terminated by ',' stored as textfile;
```

OK

Time taken: 0.035 seconds

```
hive> load data local inpath '/home/hduser/empdbnew.txt' into table empdbnew;
```

Loading data to table dsbda.empdbnew

OK

Time taken: 0.105 seconds

```
hive> select * from empdbnew;
```

OK

1	sachin	5000
2	yuvraj	3000
3	virat	1000
4	rahul	4000
5	saurav	3500
6	scott	2500
7	smith	1500
8	mathew	2400
9	Laxman	4500
10	ricky	1300
11	ram	2000

Time taken: 0.066 seconds, Fetched: 11 row(s)

```
hive> describe empinfo;
```

OK

empno	int
-------	-----

empgrade	string
----------	--------

Time taken: 0.032 seconds, Fetched: 2 row(s)

```
hive> describe empdbnew;
```

OK

eno	int
-----	-----

ename	string
-------	--------

esal	int
------	-----

Time taken: 0.026 seconds, Fetched: 3 row(s)

```
hive> select eno,ename,esal,empno,empgrade from empdbnew join empinfo
```

```
> on eno=empno;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query ID = hduser_20230510143339_a70043e2-d5d3-4794-be06-88357b16159c

Total jobs = 1

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/hduser/apache-hive-2.3.9-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

2023-05-10 14:33:43 Starting to launch local task to process map join; maximum memory = 477626368

2023-05-10 14:33:43 Dump the side-table for tag: 1 with group count: 11 into file: file:/tmp/hduser/d0c4ca8b-413b-4d29-bbb5-40bb55494a55/

hive_2023-05-10_14-33-39_907_5005826098719851142-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01--.hashtable

2023-05-10 14:33:43 Uploaded 1 File to: file:/tmp/hduser/d0c4ca8b-413b-4d29-bbb5-40bb55494a55/hive_2023-05-10_14-33-39_907_5005826098719851142-1/-local-10004/

HashTable-Stage-3/MapJoin-mapfile01--.hashtable (491 bytes)

2023-05-10 14:33:43 End of local task; Time Taken: 0.542 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1683703824975_0001, Tracking URL = [http://dpl12:8088/proxy/](http://dpl12:8088/proxy/application_1683703824975_0001/)

application_1683703824975_0001/

Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1683703824975_0001

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2023-05-10 14:33:48,986 Stage-3 map = 0%, reduce = 0%

```

2023-05-10 14:33:53,083 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.42 sec
MapReduce Total cumulative CPU time: 1 seconds 420 msec
Ended Job = job_1683703824975_0001
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.42 sec HDFS Read: 6584 HDFS Write: 413 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 420 msec
OK
1      sachin  5000    1      A
2      yuvraj  3000    2      B
3      virat   1000    3      B
4      rahul   4000    4      A
5      saurav  3500    5      B
6      scott   2500    6      B
7      smith   1500    7      D
8      mathew  2400    8      E
9      Laxman  4500    9      C
10     ricky   1300   10      C
11     ram     2000   11      D
Time taken: 14.218 seconds, Fetched: 11 row(s)
hive> create table flight(id int, Source string, Dest string, at string, dt string, delay
int)
    > row format delimited fields terminated by "," stored as textfile location
    > "/home/hduser/Desktop/flight";
OK
Time taken: 0.037 seconds
hive> load data local inpath '/home/hduser/flight.txt' into table flight;
Loading data to table dsbda.flight
OK
Time taken: 0.516 seconds
hive> select * from flight;
OK
1      Pune    Mumbai  10:25am 11:25am 5
2      Pune    Kolkata  7:00am  7:30am  2
3      Mumbai  Pune    12:30pm 12:45pm 1
4      Mumbai  Delhi   2:00pm  2:45pm 10
Time taken: 0.049 seconds, Fetched: 4 row(s)
hive> create index flight_index
    > on table flight(delay)
    > AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'
    > WITH DEFERRED REBUILD;
OK
Time taken: 0.322 seconds
hive> show index on flight;
OK
flight_index      flight      delay
dsbda__flight_flight_index__ compact
Time taken: 0.11 seconds, Fetched: 1 row(s)
hive> select avg(delay) from flight;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future
versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X
releases.
Query ID = hduser_20230510143920_33d5e238-7e3b-4dbc-8cf3-de775ce9e09b
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
    set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
    set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
    set mapreduce.job.reduces=<number>
Starting Job = job_1683703824975_0002, Tracking URL = http://dpl12:8088/proxy/
application_1683703824975_0002/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1683703824975_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-05-10 14:39:24,311 Stage-1 map = 0%, reduce = 0%

```

2023-05-10 14:39:28,396 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.94 sec
2023-05-10 14:39:31,446 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 1.95 sec
MapReduce Total cumulative CPU time: 1 seconds 950 msec
Ended Job = job_1683703824975_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.95 sec HDFS Read: 8682 HDFS Write:
103 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 950 msec
OK
4.5
Time taken: 11.777 seconds, Fetched: 1 row(s)